

Contents

1	Basic Terminologies	3
1.1	Population	3
1.2	Sample	3
1.3	Histogram	3
1.4	Law of large numbers	4
1.4.1	Weak Law of large numbers	4
1.4.2	Strong Law of large numbers	4
2	Measure of Central Tendency	5
2.1	Mean/Expected Value	5
2.2	Median	5
2.3	Mode	5
3	Measure of Spread	5
3.1	Variance	5
3.2	Standard Deviation	6
3.3	Inter Quartile Range(IQR)	6
4	Probability Distributions	6
4.1	Discrete Distributions	6
4.2	Continuous Distributions	7
5	Discrete Probability Distributions	7
5.1	Benoulli Distribution	7
5.1.1	Probability Mass Function(PMF)	7
5.1.2	Statistical Parameters	7
5.1.3	Examples	8
5.2	Binomial	8
5.2.1	Probability Mass Function(PMF)	8
5.2.2	Statistical Parameters	9
5.2.3	Examples	9
5.3	Negative-Binomial	10
5.4	Multinomial	10
5.5	Geometric	10
5.6	Hypergeometric	10
5.7	Poisson	10

5.8	Discrete Uniform	10
5.8.1	Use cases	10
6	Shape fo the Distributions	10
6.1	Skewness	10
6.1.1	Right-skewed	10
6.1.2	Left-skewed	10
7	Evaluation Metrics for Regression	10
7.1	Mean Absolute Error	10
7.1.1	Gradient of MAE	11
7.2	Mean Squared Error	11
7.2.1	Gradient of MSE	11
7.3	Root Mean Squared Error(RMSE)	11
7.3.1	Gradient of RMSE	12
7.4	R-Square(R^2)	12
7.5	Adjusted R^2	13
7.6	Mean Absolute Percentage Error(MAPE)	13
7.7	Huber Loss	14
7.7.1	Gradient of Huberloss	14
8	Evaluation Metrics for Classification	14
8.1	Basic Terminologies	14
8.2	Accuracy	15
8.3	Precision	15
8.4	Negative Predicted Value (NPV)	15
8.5	Recall	16
8.6	Specificity	16
8.7	F-score	16
8.7.1	Derivation	16
9	AUC-ROC curve	17
10	Regularization	17
10.1	17
11	Hypothesis Testing	17
12	TODO	17

Statistics For Data Science

Akash Tesla

July 2025

1 Basic Terminologies

1.1 Population

An entire set of items you want to study

1.2 Sample

A subset of population used to estimate statistical behavior of the whole population

1.3 Histogram

A histogram is a graphical representation of numerical data that groups the data into bins and displays the frequency of data points within each bin as bars




Figure 1: Example of a Histogram

1.4 Law of large numbers

As the number of trials (or samples) increases, the sample average (or empirical mean) will converge to the expected value (or population mean).

1.4.1 Weak Law of large numbers

The weak law states that the sample average of a sequence of independent identically distributed(i.i.d.) random variables converges in probability to the expected value as the number of samples goes to infinity

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i \xrightarrow{p} \mu \quad \text{as } n \rightarrow \infty$$

which means,

$$\forall \varepsilon > 0, \lim_{n \rightarrow \infty} \mathbf{p}(|\bar{X}_n - \mu| > \varepsilon) = 0$$

1.4.2 Strong Law of large numbers

The strong law states that the sample average of a sequence of i.i.d. random variables converges almost surely to the expected value as the number of samples goes to infinity

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i \xrightarrow{a.s.} \mu \quad as \ n \rightarrow \infty$$

Which means,

$$P(\lim_{n \rightarrow \infty} \bar{X}_n = \mu) = 1$$

2 Measure of Central Tendency

2.1 Mean/Expected Value

Average of all data points, sensitive to outliers since a single large outlier could easily skew mean

$$\mu = \frac{\sum x_i}{n}$$

2.2 Median

The middle data point when data are stored, robust to outliers

2.3 Mode

The most frequent data point of the dataset

3 Measure of Spread

Range: Difference between minimum value and maximum value

$$Range = x_{max} - x_{min}$$

3.1 Variance

Average squared deviation

$$\begin{aligned}\sigma^2 &= E[(X - \mu)^2] \\ \sigma^2 &= E[(X - E[X])^2] \\ \sigma^2 &= E[X^2] - (E[X])^2\end{aligned}$$

$$\sigma^2 = \frac{\sum (x_i - \mu)^2}{n} (Population)$$

$$s^2 = \frac{\sum (\bar{x}_i - \mu)^2}{n - 1} (Sample)$$

3.2 Standard Deviation

Root of Variance

$$\sigma = \sqrt{\sigma^2}$$

3.3 Inter Quartile Range(IQR)

Difference between 75th Percentile/3rd Quartile and 25th Percentile/1st Quartile, it is used for outlier detection

$$IQR = Q_3 - Q_1$$

We calculate lower bounds and upper bounds to detect the outliers

$$\text{lower bound} = Q_1 - 1.5 \times IQR$$

$$\text{upper bound} = Q_3 + 1.5 \times IQR$$

the data points which values outside of the bounds is considered to be outliers, for more extreme detection $3 \times IQR$ is also used

4 Probability Distributions

4.1 Discrete Distributions

A discrete probability distribution describes the probability of occurrence of each value of a discrete random variable

- Discrete random variable: Countable values like 1,2,3
- Each individual value has an associated probability
- The sum of probabilities for all possible values is 1

$$\sum_i P(X = x_i) = 1$$

4.2 Continuous Distributions

5 Discrete Probability Distributions

5.1 Benoulli Distribution

The benouli distribution is a discrete probability distribution for a random variable which takes only two possibilities, Sucess or a failure

5.1.1 Probability Mass Function(PMF)

$$P(X = x) = \begin{cases} p & \text{if } x=1 \\ 1 - p & \text{if } x=0 \\ 0 & \text{Otherwise} \end{cases}$$

Also written as

$$P(X = x) = p^x(1 - p)^{1-x}, \quad \text{for } x \in \{0, 1\}$$

5.1.2 Statistical Parameters

Mean

Mean is the expected value over many repetitions of the same single-trial experiment, thus it would be p since, p is probability of 1 appearing and $(1-p)$ is probability of 0 appearing

$$\mu = 1 \times (p) + 0 \times (1 - p)$$

$$\mu = p$$

Variance

Variance can be defined as $\sigma^2 = E(X^2) - (E(x))^2$, Refer Variance chapter. For Bernoulli distribution, $E(X^2) = p$, $E(X) = p$, substituting we get

$$\sigma^2 = p - p^2$$

$$\sigma^2 = p(1 - p)$$

Mode

Mode for Bernoulli would what ever the outcome which is more favored, which can be defined as

$$Mode = \begin{cases} 1 & \text{If } p > 0.5 \\ 0 & \text{If } p < 0.5 \end{cases}$$

5.1.3 Examples

- Will it rain tomorrow?
- Will this patient recover?
- Will this product be defective?

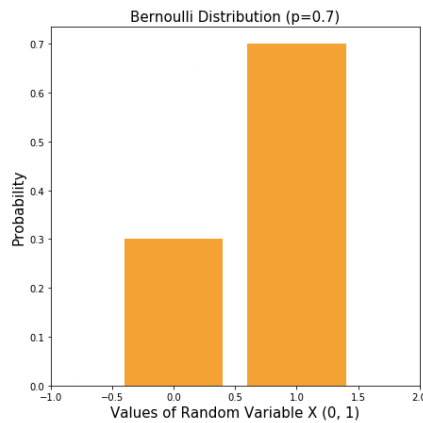


Figure 2: Example of a Bernoulli Distribution

5.2 Binomial

Binomial Distribution is a discrete probability distribution that models the probability of obtaining a specific number of successes in a fixed number of independent trials(n), these independent trials are just Bernoulli trials, you could see the similarity between them in statistical parameters

5.2.1 Probability Mass Function(PMF)

$$P(X = x) = nCx \times p^x \times (1 - p)^{(n-x)}$$

where,

n - no of trials,

p - probability of success

x - number of success

5.2.2 Statistical Parameters

Mean

Mean represents Average number of success from your trials which would be number of trials (n) times probability of success (p)

$$\mu = n \times p$$

Variance

Variance represents Expected variance between mean and data points

$$\sigma^2 = n \times p \times (1 - p)$$

Mode

$$Mode = \begin{cases} \text{floor}(n+1)p & \text{if } (n+1)p \text{ is not an Integer} \\ \text{floor}((n+1)p), \text{ floor}((n+1)(1-p)) & \text{if } (n+1)p \text{ is an Integer} \end{cases}$$

$$\text{Mode(if } p = 0.5) = \begin{cases} \frac{n}{2} & \text{if } (n+1)p \text{ is not an Integer} \\ \frac{(n-1)}{2}, \frac{(n+1)}{2} & \text{if } (n+1)p \text{ is an Integer} \end{cases}$$

5.2.3 Examples

- How many patients will recover out of 50?
- How many rainy days this month?
- How many defective products in a batch of 1000?

5.3 Negative-Binomial

5.4 Multinomial

5.5 Geometric

5.6 Hypergeometric

5.7 Poisson

5.8 Discrete Uniform

5.8.1 Use cases

- When there is only one trial
- When the outcome is binary True/False Yes/No

6 Shape fo the Distributions

6.1 Skewness

Measure of Asymmetry

6.1.1 Right-skewed

tail on the right ($mean > median$)

6.1.2 Left-skewed

tail on the left ($mean < median$)

7 Evaluation Metrics for Regression

7.1 Mean Absolute Error

$$MAE = \frac{1}{n} \sum |y_i - \hat{y}_i|$$

- Robust to outliers, treats all errors equally doesn't square the errors like RMSE,MSE..etc

- It's used when your model can tolerate moderate outliers
- Interpretability - Has same unit as the thing you are predicting/easy to understand
- Gives out constant gradient (bad for gradient based loss function)

7.1.1 Gradient of MAE

$$\frac{d}{d\hat{y}}|y - \hat{y}| = \begin{cases} +1 & \text{if } \hat{y} < y \\ -1 & \text{if } \hat{y} > y \\ \text{undefined} & \text{if } \hat{y} = y \end{cases}$$

As you can see no matter how far the error is from true value it always gives a constant gradient as it treats every error as same stics

7.2 Mean Squared Error

$$MSE = \frac{1}{n} \sum (y_i - \hat{y}_i)^2$$

- Penalizes large errors/outliers
- Gives out strong gradient signals

7.2.1 Gradient of MSE

$$\frac{dMSE}{d\hat{y}} = -\frac{2}{n}(y - \hat{y})$$

It points in the direction of the error, and it grows linearly with size of the error Larger the gradient, when prediction are more wrong → model adjusts faster

7.3 Root Mean Squared Error(RMSE)

$$RMSE = \sqrt{MSE}$$

- It combines interpretability of MAE and sensitive to errors of MSE
- It has smooth gradient curves just like MSE, and it's preferred for gradient descent

7.3.1 Gradient of RMSE

$$\frac{dRMSE}{\hat{y}_i} = \frac{1}{n \times RMSE}(\hat{y}_i - y_i)$$

1. The gradient strength changes with RMSE, if your RMSE is very large the gradient becomes small, and if your RMSE is very small the gradient becomes large.
2. It makes RMSE a Non-constantly scaled loss
3. MSE is preferred over RMSE in training, but RMSE is preferred while reporting for interpretability

7.4 R-Square(R^2)

R^2 is the coefficient of determination. it tells how well your regression model explains the variation in the dependent variable(Y) using independent variables(X)

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}}$$

where,

- $SS_{res} = \sum(y_i - \hat{y}_i)^2 \rightarrow$ Residual sum of squares(error)/MSE
- $SS_{tot} = \sum(y_i - \bar{y}_i)^2 \rightarrow$ Total sum of squares (total variability)

$$R^2 = 1 - \frac{\sum(y_i - \hat{y}_i)^2}{\sum(y_i - \bar{y}_i)^2}$$

Or, it can also be written more intuitively as

$$R^2 = 1 - \frac{MSE}{\sigma^2}$$

- Let us understand the formula (1-) operator just switches from maximizing to minimizing so you can ignore that.
- $\frac{MSE}{\sigma^2}$ Explains how well our model performs to a model that just predicts mean everytime, so if the ratio is 1, then our model is same as the dumb model, we have to reduce the ratio but the world likes "more the better" approach add (1-) operator we have to maximize the error and it's called as R^2

- R^2 ranges from $(-\infty, 1]$

7.5 Adjusted R^2

$$R_{adj}^2 = 1 - \left(\frac{(1 - R^2)(n - 1)}{n - k - 1} \right)$$

The above mentioned is textbook formula but we use our simplified representation for R^2

$$R^2 = 1 - \frac{MSE}{\sigma^2}$$

so, R_{adj}^2 would be

$$R_{adj}^2 = 1 - \frac{MSE_{adj}}{\sigma_{adj}^2}$$

MSE adjusted accounts for the number of freedoms used up to predict the data, which is K, represents number of parameters like number of predictors, number of bias

$$MSE_{adj} = \frac{\sum (y_i - \hat{y}_i)^2}{n - k}$$

Variance adjusted for number of freedoms used up which is 1 (mean), thus it'd be n-1 insted of n

$$\sigma_{adj}^2 = \frac{\sum (y_i - \mu)^2}{n - 1}$$

Substituting we get,

$$R_{adj}^2 = 1 - \left(\frac{MSE}{\sigma^2} \times \frac{n - 1}{n - k} \right)$$

where

- n - number of samples/ training samples
- k - number of parameters

7.6 Mean Absolute Percentage Error(MAPE)

MAPE is a metric used to measure accuracy of a predictive model. It expresses the prediction error as the percentage of actual values

$$MAPE = \frac{1}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right| \times 100$$

MAPE is just like MAE but it gives out the error in percentage thus it's easier to interpret

7.7 Huber Loss

Huber loss is a robust loss function/evaluation metric that has both strengths of MAE and MSE

$$L_{\delta} = \begin{cases} \frac{1}{2}(y - \hat{y})^2 & \text{for } |y - \hat{y}| \leq \delta \\ \delta \cdot (|y - \hat{y}| - 1/2\delta) & \text{otherwise} \end{cases}$$

where,

- δ is a hyperparameter that controls the behavior between MSE and MAE behavior

7.7.1 Gradient of Huberloss

$$\frac{\partial L}{\partial \hat{y}} = \begin{cases} -(y - \hat{y}) & \text{if } |y - \hat{y}| \leq \delta \\ -\delta \cdot \text{sign}(y - \hat{y}) & \text{otherwise} \end{cases}$$

example graph goes here

8 Evaluation Metrics for Classification

8.1 Basic Terminologies

True Positive

Correctly predicted positive class

False Positive

Falsely Predicted Positive class, actually negative

True Negative

Correctly predicted negative class

False Negative

Falsely predicted negative class, actually positive

8.2 Accuracy

Out of all predictions how many are correct, ratio between correct predictions and total predictions is accuracy

$$Accuracy = \frac{\text{Correct predictions}}{\text{Total predictions}}$$

can also be written as,

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN}$$

8.3 Precision

Out of all my positive class prediction how much did I get correctly, ratio between correct positive class prediction and total positive class predicted

$$\text{Precision} = \frac{\text{No of correct positive class predicted}}{\text{No of positive class predicted}}$$

can also be written as,

$$\text{Precision} = \frac{TP}{TP + FP}$$

Use when FP is costly, when you want every positive prediction to be trust worthy, measures reliability

8.4 Negative Predicted Value (NPV)

Out of all my negative predictions how much did I get correctly, ratio between

$$\text{NPV} = \frac{\text{No of correct negative class predicted}}{\text{No of negative class predicted}}$$

can also be written as,

$$\text{NPV} = \frac{TN}{TN + FN}$$

Use when FN is costly, when you want every negative class to be trust worthy

8.5 Recall

Out of all positive cases how many did I predict correctly, ratio of correctly predicted positive class and total positive cases

$$\text{Recall} = \frac{\text{No of correctly predicted positive class}}{\text{No of actual positive class}}$$

$$\text{Recall} = \frac{TP}{TP + FN}$$

Use when FN is costly/ wrongly classifying , did it recall everything, predict all positive cases

8.6 Specificity

Out of all negative cases how many did I predict correctly, Basically Recall for negative class

$$\text{Specificity} = \frac{\text{No of correctly predicted negative class}}{\text{No of actual negative class}}$$

$$\text{Specificity} = \frac{TN}{TN + FP}$$

Use when FP is costly, when predicting negative class is important

8.7 F-score

It's a harmonic mean between precision and recall

$$F_\beta = \frac{(1 + \beta^2) \cdot \text{Precision} \cdot \text{Recall}}{(\beta^2 \cdot \text{Precision}) + \text{Recall}}$$

8.7.1 Derivation

$$\frac{1}{F_\beta} = \frac{w_p}{\text{Precision}} + \frac{w_r}{\text{Recall}}$$

We want harmonic ratio of precision and recall,

$$\frac{w_r}{w_p} = \beta^2$$

We want the weights to add up to one

$$w_r + w_p = 1$$

By solving we get

$$w_r = \frac{\beta^2}{1 + \beta^2}$$

$$w_p = \frac{1}{1 + \beta^2}$$

9 AUC-ROC curve

10 Regularization

10.1

11 Hypothesis Testing

12 TODO

1. complete the Distributions, some intro to probability
2. p-value and how to make conclusions with stats
3. Time series, ML, other stuff